



# Machine Learning

---

## Feature Learning



# Feature Learning

---

- Feature representations are important for supervised learning problems as well as Reinforcement Learning problems
  - Feature representation can change the complexity of the decision boundary
  - Feature representation can change the dimensionality of the problem
- Features can be either designed or learned
  - Supervised: e.g. ANN hidden units learn features
  - Unsupervised: e.g. cluster IDs as features



# Unsupervised Feature Learning

---

- Unsupervised features have to be built around general principles for “good” features
  - Metric for “good” features has to be built into the algorithm
- Clustering can be used to form features IDs can be used as discrete features (similar items have the same/similar features)
  - Cluster ID as discrete feature
  - Cluster probability as continuous feature
- Other criteria for “good” features can be used



# Unsupervised Feature Learning

---

- A common criterion for “good” features are
  - How precisely they can represent the data
  - How compact the basis is
- Representation accuracy is aimed at ensuring that the unsupervised feature learning does not lose significant amounts of information
  - Information loss might make some subsequent tasks impossible to do/learn
- Compactness is aimed at simplicity
  - Reduce overfitting
  - Reduce complexity for subsequent tasks



# Principal Component Analysis

- Principal component analysis (PCA) is one of the most used approaches for unsupervised learning of a compact feature space
  - Uses both accuracy of representation of data and compactness of the representation
  - Assumes linear representation of data in terms of learned, constant, unit length basis vectors

$$\hat{x}^{(i)} = \bar{x} + \hat{a}_k f_k(x^{(i)}) \hat{u}_k$$

- Accuracy of representation is defined in terms of a common criterion for “good” features are

$$E_{k,f,\hat{u},D} = \hat{a}_i \left\| x^{(i)} - \hat{x}^{(i)} \right\|^2 = \hat{a}_i \left\| x^{(i)} - \left( \bar{x} + \hat{a}_{j=1}^k f_j(x^{(i)}) \hat{u}_j \right) \right\|^2$$



# Principal Component Analysis

- PCA basically forms a new basis for the data in such a way that for every number of features,  $k$ , the resulting basis minimizes the square reconstruction error
  - Each feature tries to capture as much of the remaining data variation as possible, reducing the squared error as much as possible

- Feature value that minimizes the error for all  $k$ :

$$\hat{f}_k(x^{(i)}) = \hat{u}_k^T (x^{(i)} - \bar{x})$$

- Corresponding basis vectors have to minimize the error

$$\hat{u}^* = \arg \min_u \left\| x^{(i)} - \left( \bar{x} + \sum_{j=1}^k \hat{f}_j(x^{(i)}) \hat{u}_j \right) \right\|^2$$



# Principal Component Analysis

- To solve for the basis it is important to note:
  - The vectors have to be orthogonal
  - The error for  $k=d$  is 0
- Thus the error for a  $k$  is equal to the value of the higher components

$$\begin{aligned} E_{k,f,\hat{u},D} &= \hat{a}_i \left\| \hat{a}_{j=k+1}^d \hat{u}_j^T (x^{(i)} - \bar{x}) \hat{u}_j \right\|^2 \\ &= \hat{a}_i \hat{a}_{j=k+1}^d \left( \hat{u}_j^T (x^{(i)} - \bar{x}) \right)^2 \\ &= \hat{a}_{j=k+1}^d \hat{a}_i \left( \hat{u}_j^T (x^{(i)} - \bar{x}) \right)^2 = \hat{a}_{j=k+1}^d \hat{a}_i \left( (x^{(i)} - \bar{x})^T \hat{u}_j \right)^2 \\ &= \hat{a}_{j=k+1}^d \hat{a}_i \hat{u}_j^T (x^{(i)} - \bar{x}) (x^{(i)} - \bar{x})^T \hat{u}_j = \hat{a}_{j=k+1}^d \hat{u}_j^T S \hat{u}_j \end{aligned}$$



# Principal Component Analysis


---

- Solving from top down, starting with  $k=d$  we can notice that  $\hat{u}_d = \arg \min_{\hat{u}} \hat{u}_d^T S \hat{u}_d$  has its solution for the smallest eigenvector of  $\Sigma$ 
  - In the same way, each earlier basis vector corresponds to the next smaller eigenvector of  $\Sigma$
- Principal components of a data set are the eigenvectors of normalized data's covariance matrix in order of increasing eigenvalue
  - If scales of original dimensions are incompatible, data can be normalized with standard deviation



# Principal Component Analysis

- PCA example





# Principal Component Analysis

---

- PCA algorithm
  - Create  $n \times d$  data matrix  $D$
  - Normalize columns by subtracting column average
    - If desired, normalize columns with standard deviation
  - Compute scaled covariance matrix  $\Sigma = D^T D$  of data
  - Find eigenvectors and eigenvalues of  $\Sigma$
  - Sort by eigenvalue for Principal Components
- Eigenvalue indicates loss when not using the principal component
  - Shorter representation by ignoring higher components



# Principal Component Analysis

---

- If  $d$  is large the eigenvector calculation becomes expensive and potentially numerically unstable
  - Can solve using Singular Value Decomposition
$$D = USV^T$$
    - $S$  is a diagonal matrix of the eigenvalues of  $D^TD$
    - The columns of  $V$  are the eigenvectors of  $D^TD$
  - SVD is more stable and often more efficient

# Principal Component Analysis

- PCA is one of the most commonly used feature learning approaches
  - E.g. Eigenfaces:
    - 25 most significant principal components of a set of face images:



# Eigenfaces

- Using these 25 features we use nearest neighbor to identify the person
  - Reduction to a 25 dimensional representation
  - Recognition rate is above 80% for the test set
  - Reconstruction:



C. DeCoro



# Feature Learning

---

- Other techniques exist to learn a different feature representation
  - Independent component analysis (ICA)
    - Similar to PCA but finds most statistically independent components
      - Minimizes mutual information between components
      - Or: maximizes non-Gaussianity
    - Used to separate multiple sources of stochastic data
  - Sparse PCA
    - PCA where components can be in at most  $k$  dimensions
  - Sparse coding
    - PCA with a regularization term over feature values



# Self-Organizing Maps

---

- Self-organizing maps are neural networks that are trained in an unsupervised fashion
  - Hidden units are arranged in a k-dimensional lattice with a distance function
    - Weight vectors to a unit map into the lattice
    - Neighboring unit's weight vectors are weakly linked based on the distance function
    - Training "deforms" lattice to map onto data points
      - Topological mapping
  - Units compete for Best Matching Unit (BMU)
    - Units cooperate with BMU, updating based on distance to BMU



# Self-Organizing Maps



---

- Randomly initialize hidden unit weight vectors
- For each data point find the unit that has the weight vector most similar to the data (BMU)
  - Similarity is usually Cartesian distance
- Update the weights of all the units
$$w_{j,k} \leftarrow w_{j,k} + \alpha(t)Q_t(k, BMU)(x_j^{(i)} - w_{j,k})$$
  - $Q_t(k, l)$  is the similarity between the lattice location of nodes  $k$  and  $l$
- Repeat with next data point until iteration limit



# Self-Organizing Maps

- SOM example



J. Lampinen  
T. Kostainen



# Unsupervised Feature Learning

---

- Unsupervised feature learning finds possible feature representations based on characteristics built into the algorithm
  - PCA is the most commonly used
    - PCA can always perfectly represent the original data
  - ICA and sparse methods can find more features than in the original space
    - Features can be more expressive
    - Features can be more causal
  - SOM establishes a topological mapping onto a k-dimensional lattice
    - Can be seen as a non-linear PCA